Three tools. Different jobs.

Semaphore vs Qlik Sense / Qlik Cloud Analytics vs Data Noodler

Semaphore emphasizes brief-to-dashboard drafting and reviewed answers. **Data Noodler** emphasizes preparing local files and building data products. **Qlik** supplies an enterprise analytics platform with an associative engine. Their dashboard overlap does not establish feature parity.

Where each earns its place

Semaphore	Qlik	Data Noodler
Draft and review Describe a chart, inspect its binding and SQL, accept it into a playbook, and share summarized answers. [1,2]	**Governed exploration** Explore associations in an analytical model and distribute analytics through a managed platform. Deployment and capabilities vary by edition. [4,5]	**Prepare and publish** Import local files, transform them with SQL, build charts or grids, and publish dashboard outputs. [3]
0.3 engineering preview Compared here using the consolidated application in PR #17. Engineering preview; operational readiness still requires deployment validation. [1]	**Commercial product family** This comparison covers Qlik Sense / Qlik Cloud Analytics, not the entire Qlik Talend portfolio. [4,5]	**Repository implementation** Compared using the current master README. Described features were not independently benchmarked or exercised for this report. [3]

Practical selection

Consider a Semaphore pilot when analysts need help drafting recurring dashboards and leaders need reviewed snapshot answers on constrained Windows workstations.

Consider Data Noodler when the problem starts with CSV, Parquet, Arrow or SQLite files that need importing, cleaning, SQL transforms and reusable local outputs.

Evaluate Qlik when centralized governance, enterprise distribution and associative exploration drive the requirement. Validate the chosen edition and deployment.

What the overlap hides

Implementation differences matter more than a list of shared chart types.

Capability	Semaphore preview [1,2]	Qlik [4,5,6]	Data Noodler [3]
Authoring	Plain-English flag briefs; deterministic local rules; optional Fairlead assistance. Analyst reviews the result.	App and sheet authoring, expressions, associative exploration and augmented analytics; edition dependent.	SQL transforms, data grids and code-driven D3 Studio charts placed on a dashboard.
Data path	Sample / CSV in the browser; configured database queries via its PowerShell server; Athena code also present.	Data is organized into Qlik analytical applications / models. Source connections and deployment depend on the product.	Raw files become Arrow IPC datasets; a C# engine queries working data in memory or from dataset parts.
Selection model	Page cross-filtering and explicit, approved relationships. This is not evidence of the full Qlik associative engine.	Associative engine supports exploration across the analytical model.	Dashboard widgets query its local SQL engine. Qlik-style associative equivalence is not established.
Preparing data	Profiles sources, proposes catalogs and cleaning mappings; preview includes constrained joins and a SQL desk.	Data modeling and analytics capabilities span a broader product platform. Scope varies by edition.	Import column/type selection, named SQL queries or scripts, and persisted cleaned datasets.
Sharing	Accepted playbook entries and summarized snapshots for Ask view; HTML report export code. Refresh needs a running instance or scheduled action.	Managed analytics distribution; reporting, automation and collaboration vary across Cloud and client-managed offerings.	Publish module describes web app, PDF, PowerPoint and image outputs; CSV, Parquet and SQLite dataset export.
Qlik migration	Parses a supported subset of pasted load scripts and expressions. Unsupported syntax needs review. No full .qvf conversion is established.	The native platform for existing Qlik apps and their original expressions / models.	The reviewed README does not document a Qlik app migration path.

Reading this table: Semaphore preview features are implemented development work, not a production-readiness claim. Data Noodler descriptions are repository claims. No shared speed, scale, migration-fidelity or accuracy benchmark was run.

Where the work happens

Simplified flows show the main processing boundary, not every connector or endpoint.

SEMAPHORE / reviewed analytical answers

Brief + catalog
Browser drafts and validates chart bindings.

→

Execute + inspect
Browser sample/CSV engine or local server connectors.

→

Accept + share
Playbook definitions and aggregate snapshots.

Optional AI path: browser -> Semaphore server -> local Fairlead router -> configured model. A local route is the default; explicit remote routes can transmit prompts. Database traffic does not pass through Fairlead. [1,2]

DATA NOODLER / local data preparation

Raw files
CSV / Parquet / Arrow / SQLite input.

→

Local C# engine
Arrow working datasets; SQL transforms and queries.

→

Dashboard + export
D3 charts, data grids and published outputs.

A PowerShell server serves the browser UI and compiles the engine from C# source. Its README describes Arrow files on disk, memory-backed queries and part-based reads for larger datasets. [3]

QLIK / enterprise analytical model

Sources + model
Create or load the analytical application.

→

Associative engine
Explore selections and model relationships.

→

Managed analytics
Apps, visualizations and distribution.

Qlik Cloud is hosted by Qlik. Client-managed Qlik Sense uses customer-managed infrastructure. Cloud and client-managed capabilities are not interchangeable. [4-6]

Privacy follows the configured path

Semaphore and Data Noodler target local Windows use, but local software can still access network sources or export data. Semaphore model prompts can contain briefs, catalog details and selected business values. A Qlik privacy assessment must identify the exact hosting, source connections and controls being used.

Use the same pilot dataset

A useful evaluation checks answers, portability and operations before judging convenience.

A short evaluation that separates the choices

Check	What to compare
Answer correctness	Run identical totals, grouped measures, date comparisons, null cases and filtered views. Reconcile every result to an agreed baseline.
Analyst effort	Time import, cleanup, dashboard creation and a later change request. Record manual fixes and unsupported expressions, not just first-chart speed.
Workstation fit	Try the intended Windows machine, available drivers, memory limits, offline requirements and permitted shared folders.
Sharing and control	Check who can read the data, how refresh works, what exports contain, and which deployment / edition provides required controls.

Assessment: Data Noodler and Semaphore may complement one another: prepare files in Data Noodler, then test supported CSV outputs in Semaphore for reviewed dashboards. This is a proposed workflow, not a verified native integration. A Qlik replacement decision needs the pilot above; UI similarity alone is insufficient.

Sources and version boundaries

[1] Semaphore reconciled application / README - Screens captured from preview 093defc. See PR #17 for the current integration status.

[2] Existing Semaphore for Analysts HTML - Qlik terminology and migration material; reconciled against current README.

[3] Data Noodler README - Master snapshot e118257; documentation review only.

[4] Qlik Cloud Analytics evaluation guide - Official product scope; consulted 29 October 2026.

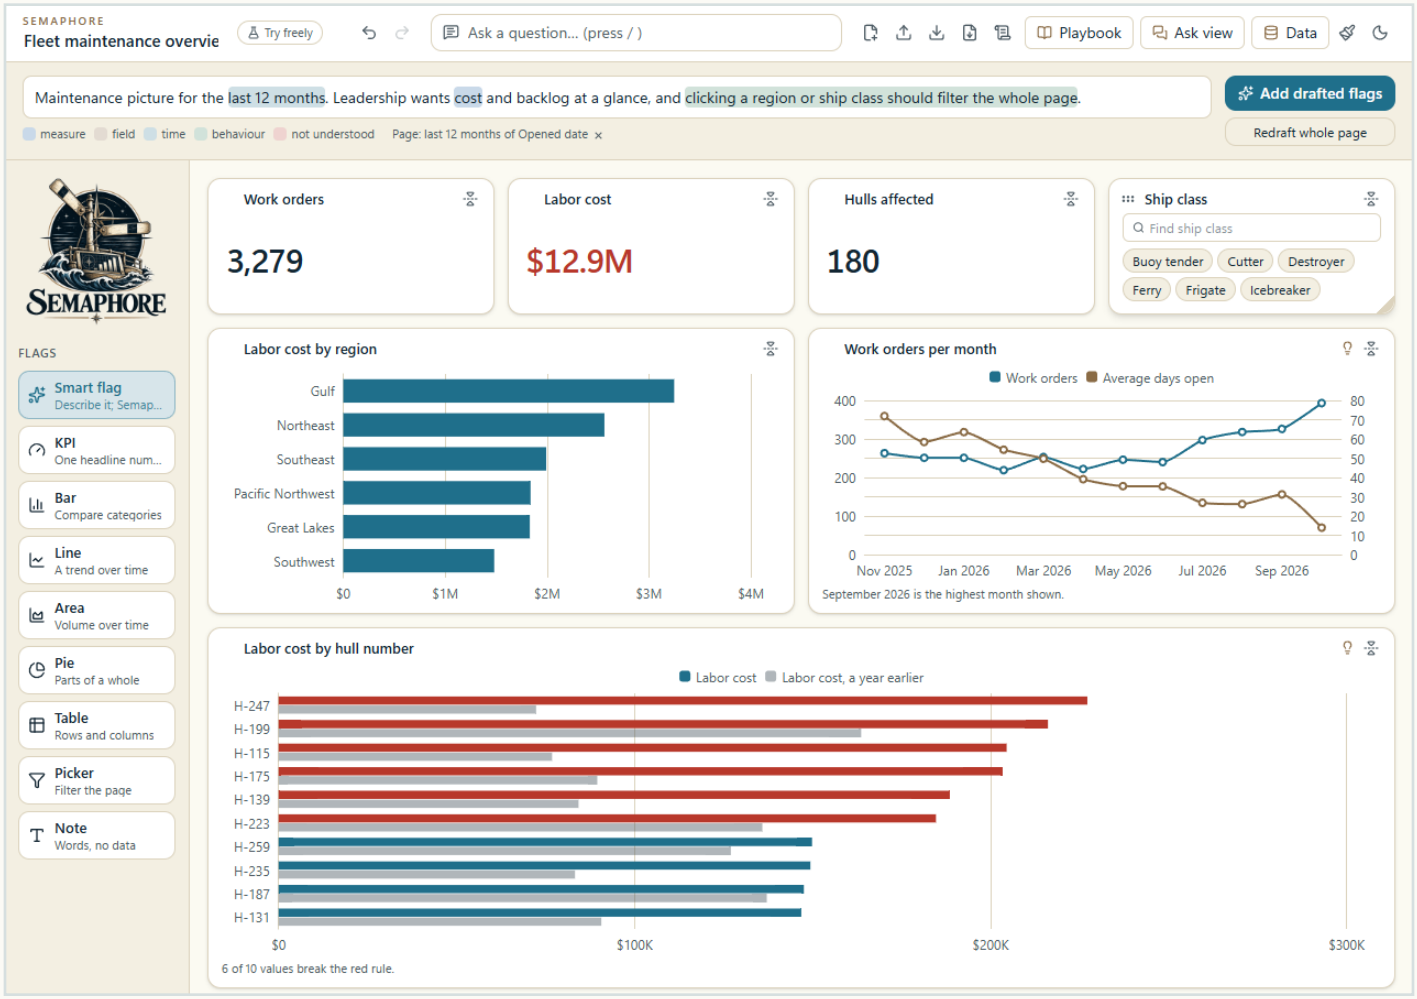
[5] Qlik Sense at a glance - Official product-family and deployment context.

[6] Qlik deployment guidance - Official Cloud versus client-managed guidance.

[7] Existing illustrated Fleet vs. enterprise briefing - Discussion links PR #74 and proposals/fleet-vs-enterprise/fleet-vs-enterprise.html. Related source, not an exact three-way comparison.

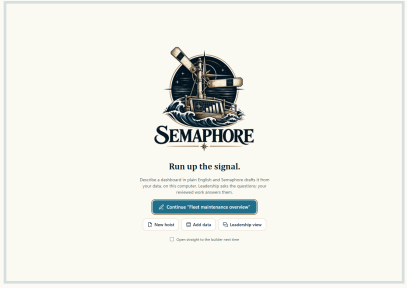
Start with the dashboard

Actual preview capture: built-in synthetic data, local rules, and a fresh browser session.



What the reader is looking at

1. Brief at the top. The analyst states the overall question and desired filtering behavior.
2. Flags in the middle. KPIs, category bars, a monthly trend and a ranked comparison share the page.
3. Tools at the edges. The palette adds chart types; Data opens the catalog; Ask view opens leadership questions.

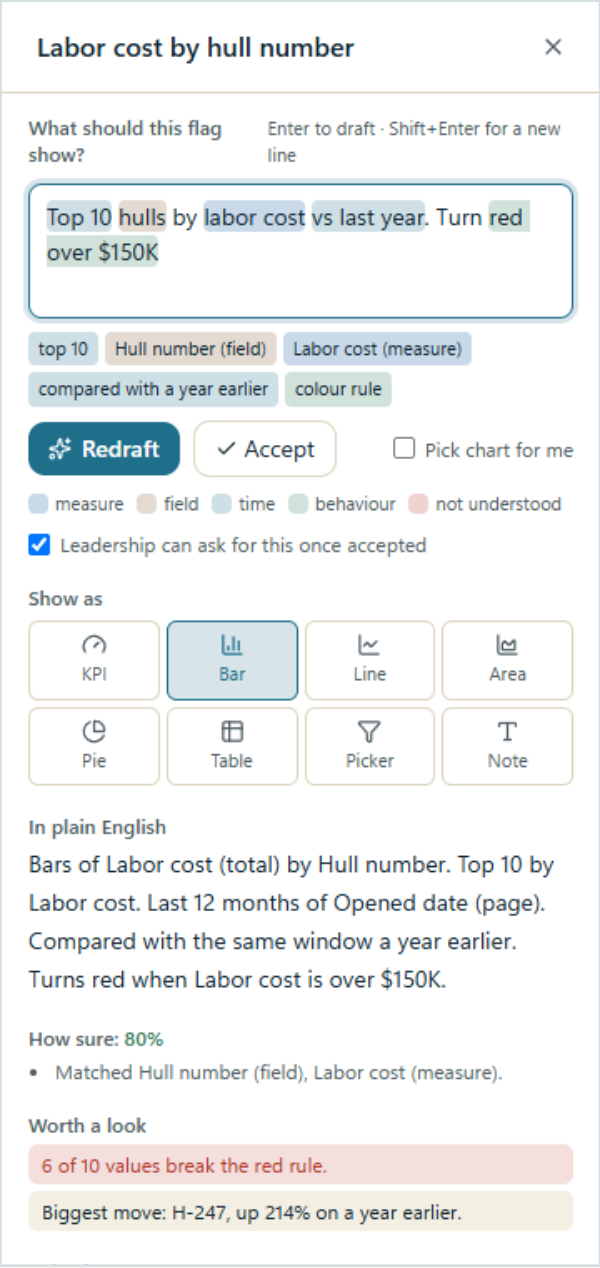


The first screen

New users can continue the example, create an empty hoist, add data or open the leadership view. This gallery follows the example into a drafted chart.

Describe, inspect, accept

The draft is a proposal the analyst can inspect and change.



Type the request
Top 10 hulls by labor cost vs last year. Turn red over $150K.

Read the interpretation
Colored annotations identify the ranking, measure, field, comparison and color rule. These make a misunderstanding easier to spot.

Check the proposed display
The inspector exposes chart kind and a plain-English readback. In this sample, it identifies a bar chart and the date window inherited from the page.

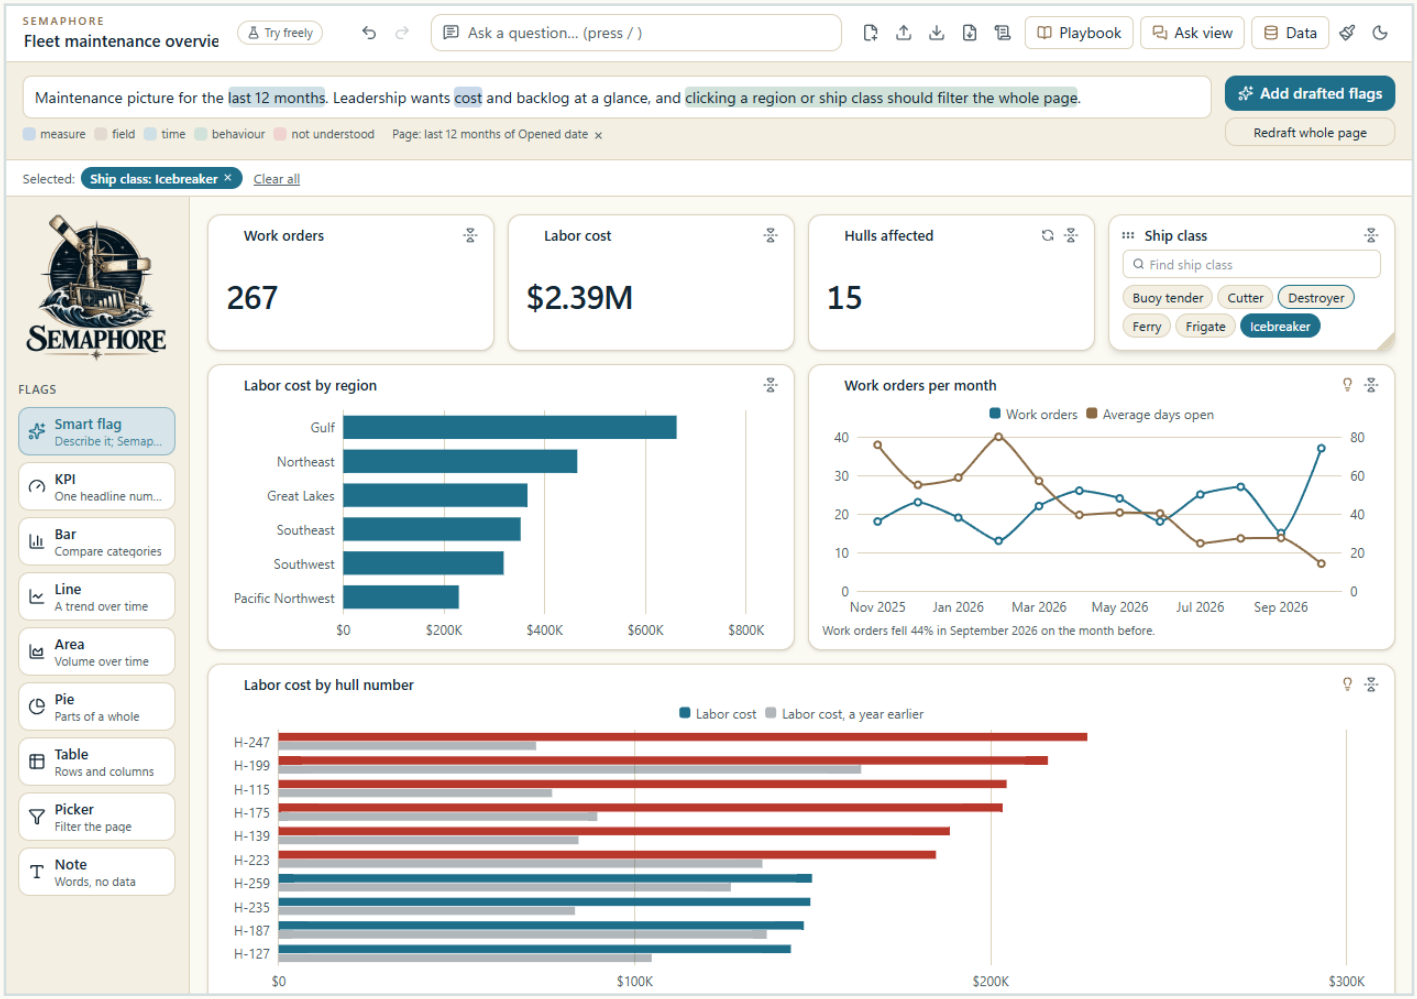
Keep the review step
The 80% indicator is the draft confidence shown by the app, not a measured probability of correctness. The analyst still checks the data and binding before accepting.

Accept useful work
Acceptance changes the flag state. Leadership publication additionally depends on the playbook and snapshot path; accepting a chart in this browser-only demo is not proof of a published shared answer.

Capture: current preview UI. Local deterministic drafting was used; no remote model or real database was called.

One selection, a new view

Choose a picker value and the page shows the selected slice.



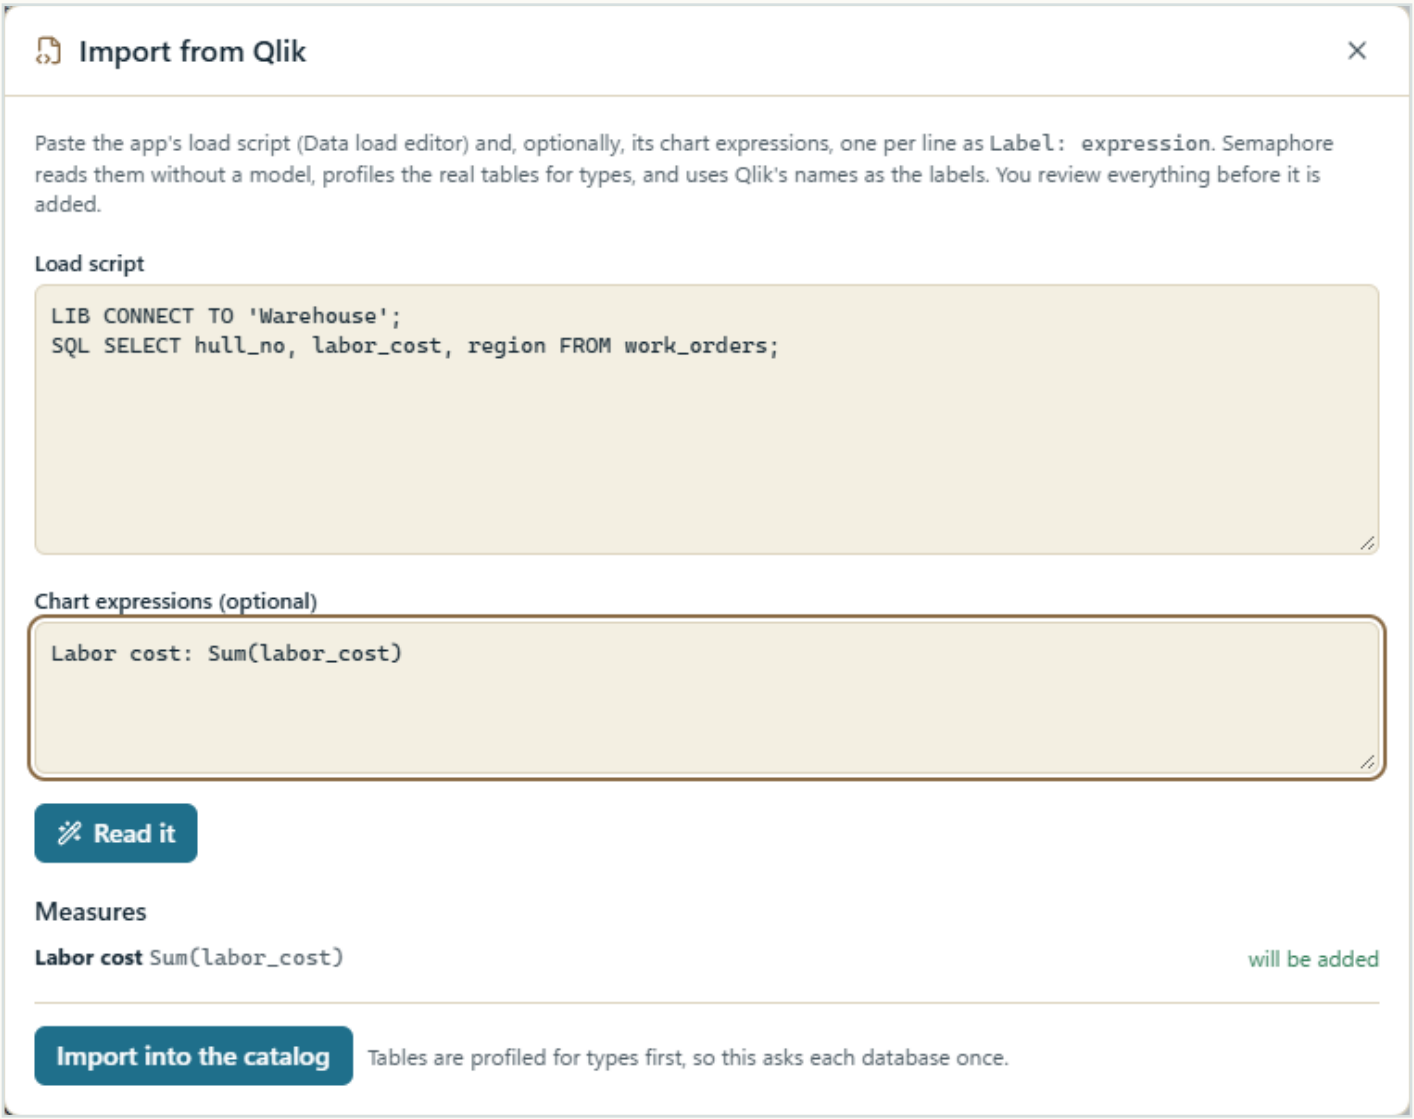
Follow the change on screen

Before selection	After choosing Icebreaker
3,279 work orders / $12.9M labor cost / 180 hulls	267 work orders / $2.39M labor cost / 15 hulls
The example covers all ship classes.	The selection chip names the active filter. The KPI values, regional bars, trend and ranked hull chart update.

These are synthetic example values from this capture, not operational measurements. Use Clear all to return to the whole page. This demonstrates Semaphore cross-filtering; it does not establish full Qlik associative-engine equivalence.

Bring over Qlik definitions

Paste supported script and expressions, read the proposed mapping, then review.



What is shown: an illustrative load script and a Sum(labor_cost) expression in the actual Import from Qlik dialog. The parser identifies a measure for review.

What was not demonstrated: a connection to Warehouse, table profiling, a completed import or a full .qvf conversion. The sample connection name is illustrative. A real source must be configured, and unsupported expressions remain review work.

Leadership asks a question

The Ask view keeps the question, answer and review status together.



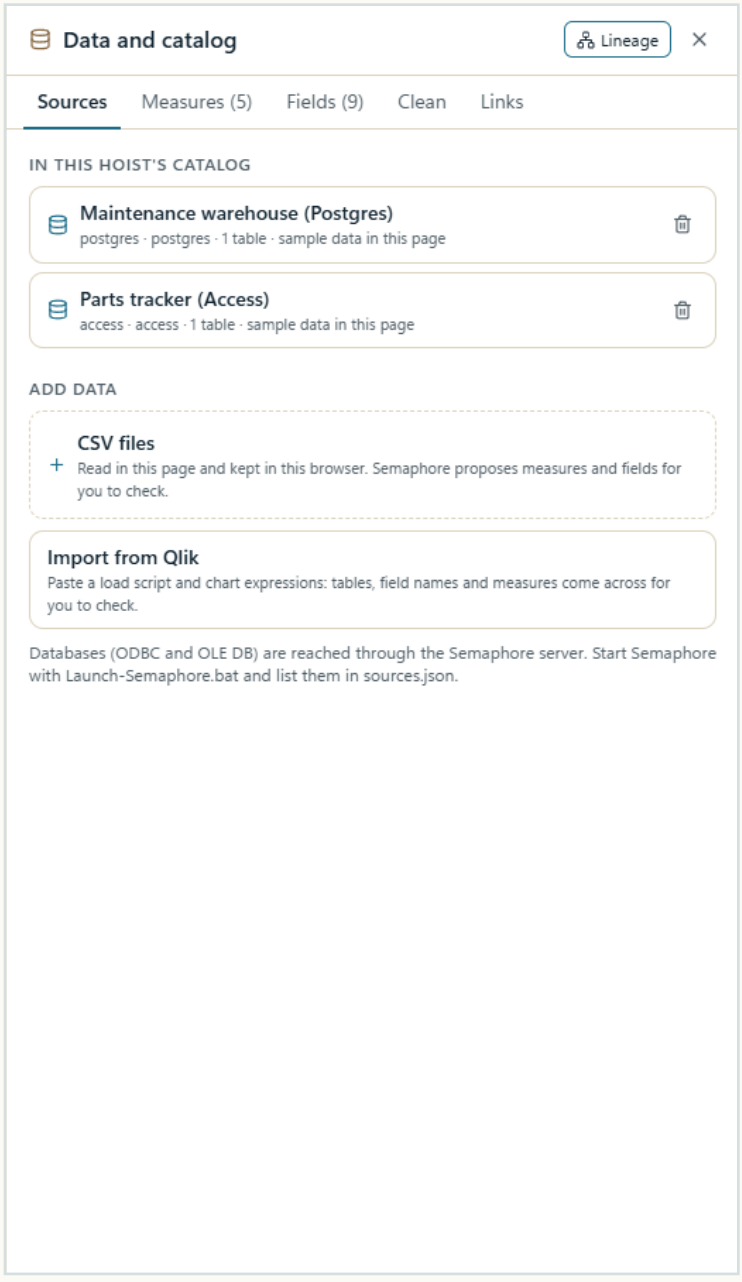
Read the status before trusting the answer

1. **Ask in ordinary language.** This example requests Labor cost by region.
2. **Read the chart and explanation.** The page shows the proposed grouping and offers follow-up questions.
3. **Check the review label.** This answer is explicitly marked Drafted by Semaphore, not reviewed.

The fresh browser-only session has 0 reviewed answers, as shown. The question is answered against synthetic example data. This capture does not demonstrate shared snapshot freshness, a server-backed analyst review queue or a live database connection.

Add data and know its source

The catalog is the bridge between familiar business terms and the data behind a chart.



Sources, measures and fields
Tabs separate source connections, analytical definitions, cleanup and relationships.

Start with local files
The CSV path reads in the browser and proposes definitions for the analyst to inspect.

Map existing Qlik work
The Qlik import entry opens the dialog on page 8. Supported definitions can be reused; migration fidelity must still be checked.

Distinguish sample from live
The source names shown here say sample data in this page. A Postgres or Access label in the example is not evidence of an active database connection.

For real databases
The panel explains that ODBC/OLE DB sources require the Semaphore server and configured source settings.

Gallery provenance: seven fresh screenshots of the built UI from feature/fairlead-intelligence (093defc), captured on 29 October 2026 with synthetic data. No screenshots were fabricated and no application state was altered to imply a successful live connection.